

# HOMEWORK 2

Answer all questions from practical assignment 2.

Section(s):

## Play With Snapshots And Images

**Question:** Open `http://195.201.14.XXX` with a browser again. What do you see and what do you think happened?

**Answer:** First we created an Apache2 web server in our container. That allowed us to visit its `localhost` and see the *Welcome* page of the properly set-up Apache. Then we created a snapshot of that container, where in our case it was called `first-snapshot`. We went about our business and altered some files, more specifically, we altered the landing page of the Apache2 web server. At this point we decided to stop our container, in order for us to restore it from the snapshot we just created. Now when we look at the *Welcome* page again, we can see that our changes were *lost*. Basically, we tested how copy-on-write works. By doing the restore, the pointer (i.e. `first-snapshot`) to our container just changed the location in pointed to. That is why we lost our changes, because when we *saved* the modified `html`, it was saved to a place where the snapshot did not point to.

---

**Question:** What do you think can go wrong when restoring a container from a snapshot?

**Answer:** Several things can go wrong in such a situation. First of all, if we do not create *regular* snapshots of our containers, we are setting up ourselves for a nasty surprise of “*If you revert to a past snapshot, all you current changes will disappear*”. But we also might put ourselves in a false sense of security, thinking that this is a valuable form of backup, but in reality it’s not. This means that even if we are keeping track of what we snapshot, we still might revert to a previous state that makes us loose valuable data. This kind-of relates to the *git-panic* situations.

---

**Question:** Can you describe what we did exactly with the above set of commands?

**Answer:** First of, we are creating a new container called `dev1`, by specifying the operating system we want the container to have. Since we’ve already done that, this time the operation is quicker. Then we issued a `destroy` command on the newly created container, thus removing any *datasets* in it. But because we *saved* our data set from the previous container, under the name `dev1.raw`, we can use the `receive` command in order to restore our file system to its previous state. Finally we have to tell `zfs` where to mount for the specific container. This is done through the `set mountpoint` command.

---

**Question:** What do you think we can use this approach for?

**Answer:** One of the uses of this approach is that we can restore a container from backup incredibly fast. We did not wait for anything, and we got our previous state as it was (given that we did have a backup ready). This also means that we can create several identical containers and jump start a task fairly quickly, without having to configure/re-configure anything.

## **Extra References**

How do snapshot backups work

Copy-on-write

What is storage snapshot technology

Crash course of ZFS

*Martin Nestorov*